

You have **2** free member-only stories left this month. [Upgrade for unlimited access.](#)

Guide to Tensorflow Object Detection (Tensorflow 2)



Rohit Prakash

Oct 5, 2020 · 8 min read ★

There are many guides out there that are very good to help you get started with setting up the TF Object Detection API, but unfortunately, most of them are written for the TF v1 API.

We will take a look at how to use the TF v2 Object Detection API to build a model for a custom dataset on a Google Colab Notebook.

Before we begin the setup, make sure to change the runtime-type in Colab to GPU so that we can make use of the free GPU provided.

1. Installing Dependencies and setting up the workspace.

Create a folder for your workspace

```
%mkdir workspace  
%cd /content/workspace
```

We will be cloning the TF repository from GitHub

```
!git clone --q https://github.com/tensorflow/models.git
```

And before we install TF Object Detection we must install Protobuf.

“The Tensorflow Object Detection API uses Protobufs to configure model and training parameters. Before the framework can be used, the Protobuf libraries must be downloaded and compiled”

```
!apt-get install -qq protobuf-compiler python-pil python-lxml
python-tk
!pip install -qq Cython contextlib2 pillow lxml matplotlib

!pip install -qq pycocotools
%cd models/research/
!protoc object_detection/protos/*.proto --python_out=.
```

Now we install the TF Object Detection API

```
%cp object_detection/packages/tf2/setup.py .
!python -m pip install .
!python object_detection/builders/model_builder_tf2_test.py
```

```
[ RUN      ] ModelBuilderTF2Test.test_invalid_faster_rcnn_batchnorm_update
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_invalid_faster_rcnn_batchnorm_update): 0.0s
I0929 14:36:41.920821 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_invalid_faster_rcnn_batchnorm_update): 0.0s
[ OK      ] ModelBuilderTF2Test.test_invalid_faster_rcnn_batchnorm_update
[ RUN      ] ModelBuilderTF2Test.test_invalid_first_stage_nms_iou_threshold
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_invalid_first_stage_nms_iou_threshold): 0.0s
I0929 14:36:41.922938 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_invalid_first_stage_nms_iou_threshold): 0.0s
[ OK      ] ModelBuilderTF2Test.test_invalid_first_stage_nms_iou_threshold
[ RUN      ] ModelBuilderTF2Test.test_invalid_model_config_proto
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_invalid_model_config_proto): 0.0s
I0929 14:36:41.923602 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_invalid_model_config_proto): 0.0s
[ OK      ] ModelBuilderTF2Test.test_invalid_model_config_proto
[ RUN      ] ModelBuilderTF2Test.test_invalid_second_stage_batch_size
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_invalid_second_stage_batch_size): 0.0s
I0929 14:36:41.925252 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_invalid_second_stage_batch_size): 0.0s
[ OK      ] ModelBuilderTF2Test.test_invalid_second_stage_batch_size
[ RUN      ] ModelBuilderTF2Test.test_session
[ SKIPPED ] ModelBuilderTF2Test.test_session
[ RUN      ] ModelBuilderTF2Test.test_unknown_faster_rcnn_feature_extractor
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_unknown_faster_rcnn_feature_extractor): 0.0s
I0929 14:36:41.926978 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_unknown_faster_rcnn_feature_extractor): 0.0s
[ OK      ] ModelBuilderTF2Test.test_unknown_faster_rcnn_feature_extractor
[ RUN      ] ModelBuilderTF2Test.test_unknown_meta_architecture
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_unknown_meta_architecture): 0.0s
I0929 14:36:41.927641 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_unknown_meta_architecture): 0.0s
[ OK      ] ModelBuilderTF2Test.test_unknown_meta_architecture
[ RUN      ] ModelBuilderTF2Test.test_unknown_ssd_feature_extractor
INFO:tensorflow:time(__main__.ModelBuilderTF2Test.test_unknown_ssd_feature_extractor): 0.0s
I0929 14:36:41.929182 139723948119936 test_util.py:1973] time(__main__.ModelBuilderTF2Test.test_unknown_ssd_feature_extractor): 0.0s
[ OK      ] ModelBuilderTF2Test.test_unknown_ssd_feature_extractor
-----
Ran 20 tests in 41.672s

OK (skipped=1)
```

The output should be similar to this.

2. Preparing the Dataset

There are two ways to go about this:

- Use a Public Labelled Dataset
- Create a Custom Labelled Dataset

You can find Public Labelled Datasets online, which are already labeled and saved in the right format, ready to be used to train.

For this tutorial, we will be creating our own dataset from scratch.

First things first, gather the images for the dataset. I will assume this step has already been done.

Now we need to label the images. There are many popular labeling tools, we will be using LabelIMG.

To install LabelIMG, execute the following code (Do it on your local Terminal since Colab does not support GUI applications):

```
pip install labelImg
```

Launch LabelImg in the folder where your images are stored.

```
labelImg imagesdir
```

Now you can start labeling your images, for more info on how to label the images follow this [link](https://github.com/heartexlabs/labelImg) (LabelImg Repository).





Labellmg

Create a label map in notepad as follows (label_map.pbtxt) with two classes for example cars and bikes:

```
item {
  id: 1
  name: 'car'
}

item {
  id: 2
  name: 'bike'
}
```

Now for creating the TFRecord files.

We can do the following:

- Create TFRecord ourselves
- Upload the annotations to Roboflow and get the dataset in TFRecord Format.

Creating the TFRecords ourselves is a bit tedious as the XML created after annotating may sometimes vary, so for the sake of ease, I suggest using [Roboflow](#) to perform the above task. They also provide an option to perform additional Data Augmentation which will increase the size of the dataset.

For your reference, here is a sample .py script to create the TFRecords manually.

```
1  import pandas as pd
2  import numpy as np
3  import csv
4  import re
5  import cv2
6  import os
7  import glob
8  import xml.etree.ElementTree as ET
9
10 import io
```

```

11 import tensorflow as tf
12 from collections import namedtuple, OrderedDict
13
14 import shutil
15 import urllib.request
16 import tarfile
17 import argparse
18
19 # os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'    # Suppress TensorFlow logging (1)
20 import tensorflow.compat.v1 as tf
21 from PIL import Image
22 from object_detection.utils import dataset_util, label_map_util
23 from collections import namedtuple
24 xml_dir = 'images/test'
25 image_dir = 'images/test'
26
27 label_map = label_map_util.load_labelmap('annotations/label_map.pbtxt')
28 label_map_dict = label_map_util.get_label_map_dict(label_map)
29
30 output_path = 'annotations/test.record'
31
32 def xml_to_csv(path):
33     """Iterates through all .xml files (generated by labelImg) in a given directory and combines
34     them in a single Pandas dataframe.
35
36     Parameters:
37     -----
38     path : str
39         The path containing the .xml files
40     Returns
41     -----
42     Pandas DataFrame
43         The produced dataframe
44     """
45
46     xml_list = []
47     for xml_file in glob.glob(path + '/*.xml'):
48         tree = ET.parse(xml_file)
49         root = tree.getroot()
50         for member in root.findall('object'):
51             value = (root.find('filename').text,
52                     int(root.find('size')[0].text),
53                     int(root.find('size')[1].text),
54                     member[0].text,
55                     int(member[4][0].text),
56                     int(member[4][1].text),
57                     int(member[4][2].text),

```

```

58         int(member[4][3].text)
59     )
60     xml_list.append(value)
61     column_name = ['filename', 'width', 'height',
62                   'class', 'xmin', 'ymin', 'xmax', 'ymax']
63     xml_df = pd.DataFrame(xml_list, columns=column_name)
64     return xml_df
65
66
67 def class_text_to_int(row_label):
68     return label_map_dict[row_label]
69
70
71 def split(df, group):
72     data = namedtuple('data', ['filename', 'object'])
73     gb = df.groupby(group)
74     return [data(filename, gb.get_group(x)) for filename, x in zip(gb.groups.keys(), gb.groups)]
75
76
77 def create_tf_example(group, path):
78     with tf.gfile.GFile(os.path.join(path, '{}'.format(group.filename)), 'rb') as fid:
79         encoded_jpg = fid.read()
80         encoded_jpg_io = io.BytesIO(encoded_jpg)
81         image = Image.open(encoded_jpg_io)
82         width, height = image.size
83
84         filename = group.filename.encode('utf8')
85         image_format = b'jpg'
86         xmins = []
87         xmaxs = []
88         ymins = []
89         ymaxs = []
90         classes_text = []
91         classes = []
92
93         for index, row in group.object.iterrows():
94             xmins.append(row['xmin'] / width)
95             xmaxs.append(row['xmax'] / width)
96             ymins.append(row['ymin'] / height)
97             ymaxs.append(row['ymax'] / height)
98             classes_text.append(row['class'].encode('utf8'))
99             classes.append(class_text_to_int(row['class']))
100
101     tf_example = tf.train.Example(features=tf.train.Features(feature={
102         'image/height': dataset_util.int64_feature(height),
103         'image/width': dataset_util.int64_feature(width),
104         'image/filename': dataset_util.bytes_feature(filename),
105         'image/source_id': dataset_util.bytes_feature(filename),

```

```

106         'image/encoded': dataset_util.bytes_feature(encoded_jpg),
107         'image/format': dataset_util.bytes_feature(image_format),
108         'image/object/bbox/xmin': dataset_util.float_list_feature(xmins),
109         'image/object/bbox/xmax': dataset_util.float_list_feature(xmaxs),
110         'image/object/bbox/ymin': dataset_util.float_list_feature(ymins),
111         'image/object/bbox/ymax': dataset_util.float_list_feature(ymaxs),
112         'image/object/class/text': dataset_util.bytes_list_feature(classes_text),
113         'image/object/class/label': dataset_util.int64_list_feature(classes),
114     })
115     return tf_example
116
117     csv_path = None
118     def main(_):
119
120         writer = tf.python_io.TFRecordWriter(output_path)
121         path = os.path.join(image_dir)
122         examples = xml_to_csv(xml_dir)
123         grouped = split(examples, 'filename')
124         for group in grouped:
125             tf_example = create_tf_example(group, path)
126             writer.write(tf_example.SerializeToString())
127         writer.close()
128         print('Successfully created the TFRecord file: {}'.format(output_path))
129         if csv_path is not None:
130             examples.to_csv(csv_path, index=None)

```

```

xml_dir = 'images/test'
image_dir = 'images/test'
output_path = 'annotations/test.record'

```

By using Roboflow you will be provided the TFRecord files automatically.

Setting up on Colab

Create folders to store all the necessary files we have just created.

```

%mkdir annotations exported-models pre-trained-models
models/my_mobilenet # my_mobilenet folder is where our training
results will be stored

```

Now upload the newly created TFRecord files along with the images and annotations to Google Colab by clicking upload files.

You could use Google Drive to store your necessary files and importing those to Google Colab should be as simple as doing a `!cp` command.

Download Pre-Trained Model

There are many models ready to download from the [Tensorflow Model Zoo](#).

Be careful in choosing which model to use as some are not made for Object Detection. For this tutorial we will be using the following model:

SSD MobileNet V2 FPNLite 320x320.

Download it into your Colab Notebook and extract it by executing:

```
%cd pre-trained-models
!curl
"http://download.tensorflow.org/models/object_detection/tf2/20200711
/ssd_mobilenet_v2_fpnlite_320x320_coco17_tpu-8.tar.gz" --output
"ssd_mobilenet_v2_fpnlite_320x320_coco17_tpu-8.tar.gz"

model_name = 'ssd_mobilenet_v2_fpnlite_320x320_coco17_tpu-8'

model_file = model_name + '.tar.gz'

tar = tarfile.open(model_file)

tar.extractall()

tar.close()

os.remove(model_file)
```

Your directory structure should now look like this:

```
workspace/
├── models/
│   ├── community/
│   ├── official/
│   ├── orbit/
│   ├── research/
│   ├── my_mobilenet/
│   └── ...
└── annotations/
```



```
|  └─ train/  
|    └─ test/  
└─ pre-trained-model/  
└─ exported-models/
```

Editing the Configuration file

In TF Object Detection API, all the settings and required information for training the model and evaluating is situated in the pipeline.config file.

Let us take a look at it:

The most important ones we will need to change are

```
batch_size: 128  
  
fine_tune_checkpoint: "PATH_TO_BE_CONFIGURED"  
  
num_steps: 50000  
num_classes: 2  
  
fine_tune_checkpoint_type: "classification"  
  
train_input_reader {  
  label_map_path: "PATH_TO_BE_CONFIGURED"  
  tf_record_input_reader {  
    input_path: "PATH_TO_BE_CONFIGURED"  
  }  
}  
  
eval_input_reader {  
  label_map_path: "PATH_TO_BE_CONFIGURED"  
  shuffle: false  
  num_epochs: 1  
  tf_record_input_reader {  
    input_path: "PATH_TO_BE_CONFIGURED"  
  }  
}
```

batch_size is the number of batches the model will train in parallel. A suitable number to use is 8. It could be more/less depending on the computing power available.

A good suggestion given on [StackOverflow](#) is:

Max batch size= available GPU memory bytes / 4 / (size of tensors + trainable parameters)

`fine_tune_checkpoint` is the last trained checkpoint (a checkpoint is how the model is stored by Tensorflow).

If you are starting the training for the first time, set this to the pre-trained-model.

If you want to continue training on a previously trained checkpoint, set it to the respective checkpoint path. (This will continue training, building upon the features and loss instead of starting from scratch).

```
# For Fresh Training
fine_tune_checkpoint: "pre-trained-
model/ssd_mobilenet_v2_fpnlite_320x320_coco17_tpu-8/checkpoint/ckpt-
0"

# For Continuing the Training
fine_tune_checkpoint:
"exported_models/your_latest_batch/checkpoint/ckpt-0"

batch_size = 8 # Increase/Decrease this value depending on how fast
your train job runs and the availability of the Compute Resources.

num_steps: 25000 # 25000 is a good number of steps to get a good
loss.

fine_tune_checkpoint_type: "detection" # Set this to detection

train_input_reader {
  label_map_path: "annotations/label_map.pbtxt"    # Set to location
of label map
  tf_record_input_reader {
    input_path: "annotations/train.tfrecord"    # Set to location of
train TFRecord file
  }
}
# Similarly do the same for the eval input reader
eval_input_reader {
  label_map_path: "annotations/label_map.pbtxt"
  shuffle: false
  num_epochs: 1
  tf_record_input_reader {
    input_path: "annotations/test.tfrecord"
  }
}
```

After editing the config file, we need to add the TensorFlow object detection folders to the python path.

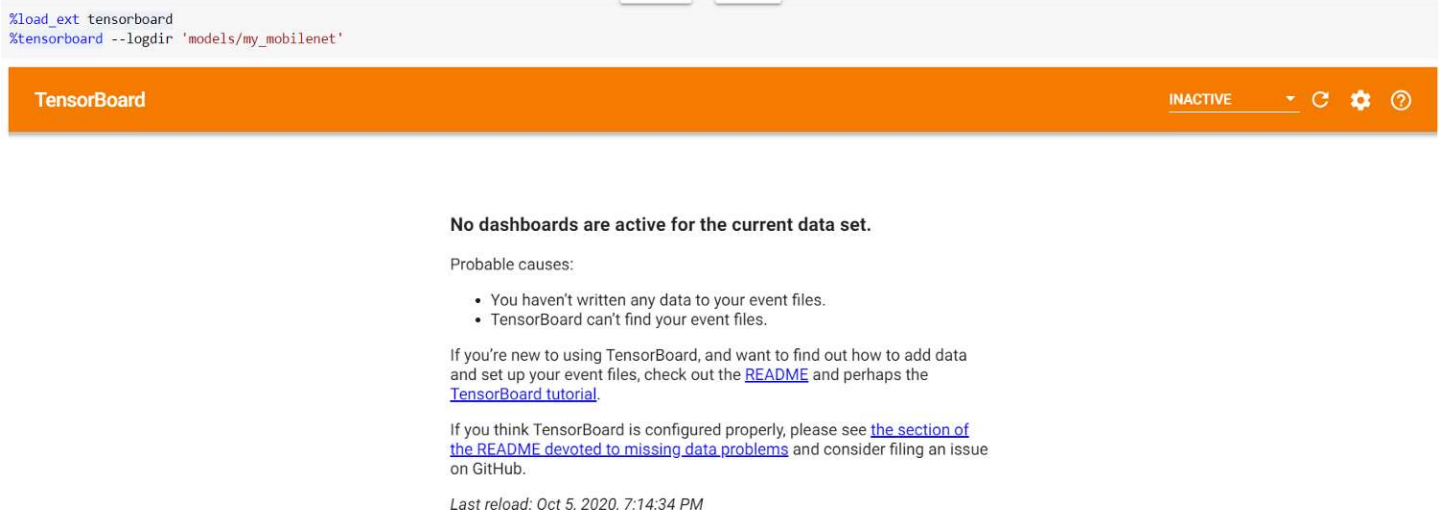
```
import os

os.environ['PYTHONPATH'] +=
':/content/window_detection/models:/content/window_detection/models
/research:/content/window_detection/models/research/slim/'
```

Setting up TensorBoard on Colab to monitor the training process

Colab has introduced inbuilt support for TensorBoard and can now be called with a simple magic command as follows

```
%load_ext tensorboard
%tensorboard --logdir 'models/my_mobilenet'
```



Cell running Tensorboard

This is how the cell will look once you execute the above command, but nothing to worry, once we start the training job, click refresh on the Tensorboard cell(Top Right) after a few minutes(The .tfevent files need to be created for us to monitor the TensorFlow logs) and you will see the output on the TensorBoard magic cell

Running the Training Job

We will copy the TensorFlow training python script to the workspace directory for ease of access.

```
!cp
'/content/window_detection/models/research/object_detection/model_main_tf2.py' .
```

The training job requires command-line arguments, namely:

- *model_dir* : This refers to the path where the training process will store the checkpoint files.
- *pipeline_config_path* : This refers to the path where the pipeline.config file is stored

Execute the following command to start the training job

```
# If you are training from scratch
!python model_main_tf2.py --model_dir=models/my_mobilenet --
pipeline_config_path=pre-trained-
model/ssd_mobilenet_v2_fpn-lite_320x320_coco17_tpu-8/pipeline.config

# Or if you are continuing from a previous training
!python model_main_tf2.py --model_dir=models/my_mobilenet --
pipeline_config_path=exported_models/pipeline.config
```

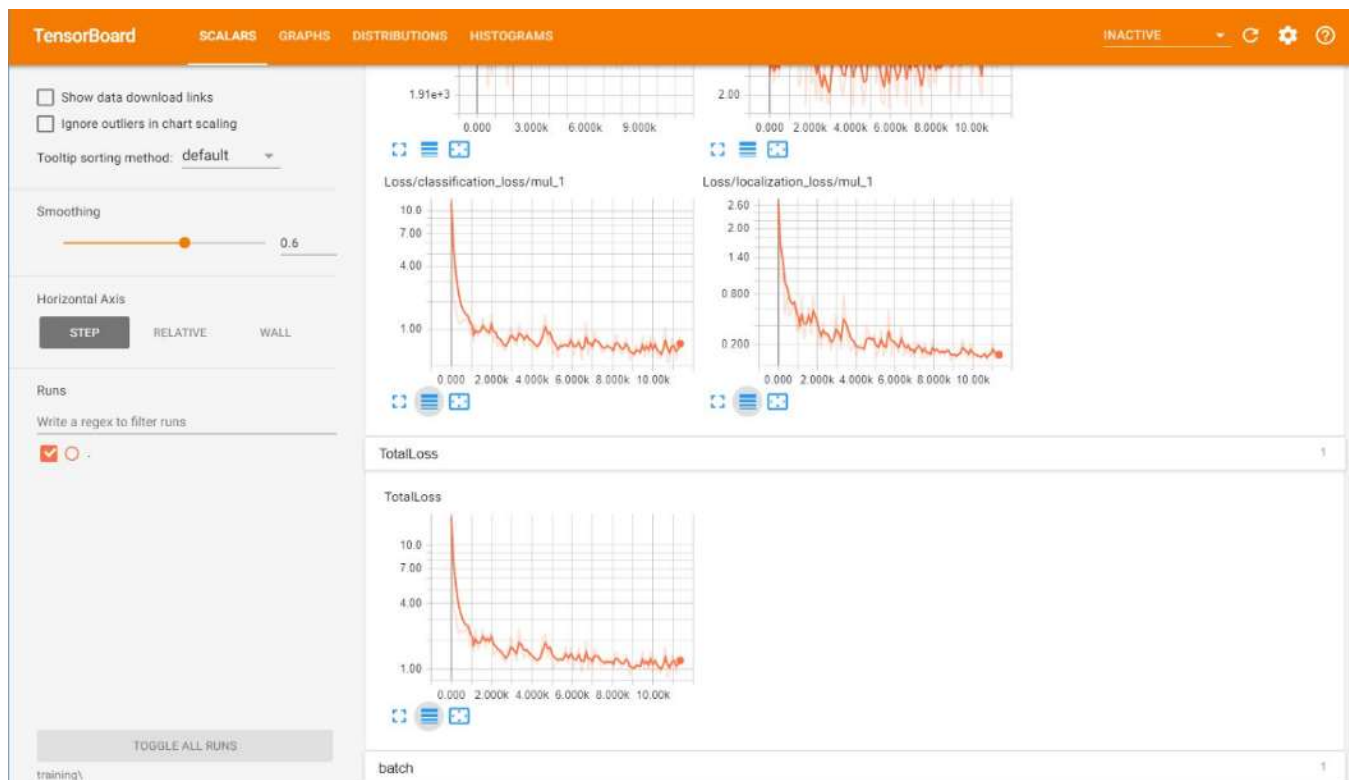
If everything goes well, the training output cell should look like this

```
I0929 19:14:03.429471 140008397502336 model_lib_v2.py:652] Step 24300 per-step time 0.355s loss=0.056
INFO:tensorflow:Step 24400 per-step time 0.335s loss=0.047
I0929 19:14:37.807487 140008397502336 model_lib_v2.py:652] Step 24400 per-step time 0.335s loss=0.047
INFO:tensorflow:Step 24500 per-step time 0.337s loss=0.056
I0929 19:15:12.283904 140008397502336 model_lib_v2.py:652] Step 24500 per-step time 0.337s loss=0.056
INFO:tensorflow:Step 24600 per-step time 0.308s loss=0.052
I0929 19:15:46.930425 140008397502336 model_lib_v2.py:652] Step 24600 per-step time 0.308s loss=0.052
INFO:tensorflow:Step 24700 per-step time 0.361s loss=0.046
I0929 19:16:21.292615 140008397502336 model_lib_v2.py:652] Step 24700 per-step time 0.361s loss=0.046
INFO:tensorflow:Step 24800 per-step time 0.338s loss=0.048
I0929 19:16:55.302360 140008397502336 model_lib_v2.py:652] Step 24800 per-step time 0.338s loss=0.048
INFO:tensorflow:Step 24900 per-step time 0.373s loss=0.062
I0929 19:17:29.647791 140008397502336 model_lib_v2.py:652] Step 24900 per-step time 0.373s loss=0.062
INFO:tensorflow:Step 25000 per-step time 0.311s loss=0.054
I0929 19:18:04.469359 140008397502336 model_lib_v2.py:652] Step 25000 per-step time 0.311s loss=0.054
```

Training Output

The output will normally update slowly. The training outputs logs only every 100 steps by default, therefore if you wait for a while, you should see a log for the loss at step 100. The speed depends on whether a GPU is being used to train and the available VRAM and many other factors, so be patient.

Refresh the TensorBoard while the training is running and you will be able to monitor the progress



Tensorboard Logs

Once the loss reaches a fairly constant value or becomes lower than 0.05 (in my case), then you can stop the training cell.

Evaluating the model

Now you can run the evaluation script to find out the mAP (Mean Average Precision) and the Loss.

Run the following in a cell:

```
!python model_main_tf2.py --model_dir=exported-models/checkpoint --
pipeline_config_path=exported-models/pipeline.config --
checkpoint_dir=models/my_mobilenet/checkpoint # The folder where the
model has saved the checkpoints during training
```

You should get an output that looks like this

```
Accumulating evaluation results...
DONE (t=0.02s).
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.833
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.911
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.859
```

Average Precision	(AP)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.800
Average Precision	(AP)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.838
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 1]	= 0.567
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 10]	= 0.873
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets=100]	= 0.894
Average Recall	(AR)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.800
Average Recall	(AR)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.896

Evaluation results

Now the evaluation script has a default timeout of 3600 seconds to wait for a new checkpoint to be generated as the script was initially intended to be running in parallel to the training job, but we are running it after the training process on Colab

You may go ahead and stop the evaluation cell from running.

Exporting the model

Now that we have our model ready, we need to save it in a format we can use it later.

We now have a bunch of checkpoints in the *models/my_mobilenet* folder. To remove all the older checkpoints and keep the latest checkpoint, I have attached a neat little python script that will do the task automatically.

```
1  output_directory = 'exported-models/'
2
3  # goes through the model is the training/ dir and gets the last one.
4  # you could choose a specific one instead of the last
5  lst = os.listdir("models/my_mobilenet/")
6  # print(lst)
7  lst = [l for l in lst if 'ckpt-' in l and '.index' not in l]
8  steps=np.array([int(re.findall('\d+', l)[0]) for l in lst])
9  last_model = lst[steps.argmax()]
10 last_model_path = os.path.join('models/my_mobilenet', last_model)
11 # print(last_model_path)
```

GetLatestCheckpoint.py hosted with ❤ by GitHub

[view raw](#)

Now to export the model, we run the export script provided by TF2, as follows:

```
!python
/content/workspace/models/research/object_detection/exporter_main_v2
.py --input_type=image_tensor \
--pipeline_config_path=pre-trained-
model/ssd_mobilenet_v2_fpnlite_320x320_coco17_tpu-8/pipeline.config
--output_directory=exported_models \
--trained_checkpoint_dir=models/my_mobilenet
```

The export script will save the model in the *exported_models* folder with the following structure

```
workspace/  
├─ exported_models/  
│   ├── checkpoint/  
│   ├── saved_model/  
│   └─ pipeline.config
```

You can now upload this folder to Google Drive or download it to save it for future use.

Inference on the model

The final step, the step that fills you with a sense of accomplishment, in this step we will test our model on a random input image and see the model predict the type of object and give its bounding box.

The entire process is a little tedious but I will attach a script that will let you perform inference directly on Google Colab

```
1  import numpy as np  
2  from PIL import Image  
3  from google.colab.patches import cv2_imshow  
4  
5  def load_image_into_numpy_array(path):  
6      """Load an image from file into a numpy array.  
7  
8      Puts image into numpy array to feed into tensorflow graph.  
9      Note that by convention we put it into a numpy array with shape  
10     (height, width, channels), where channels=3 for RGB.  
11  
12     Args:  
13         path: the file path to the image  
14  
15     Returns:  
16         uint8 numpy array with shape (img_height, img_width, 3)  
17     """  
18     return np.array(Image.open(path))  
19  
20  image_path = "PATH TO YOUR INFERENCE IMAGE"  
21  print('Running inference for {}... '.format(image_path), end='')  
22
```




```

23 image_np = load_image_into_numpy_array(image_path)
24
25 # Things to try:
26 # Flip horizontally
27 # image_np = np.fliplr(image_np).copy()
28
29 # Convert image to grayscale, (You could uncomment this to try and see how the model reacts to
30 # image_np = np.tile(
31 #     np.mean(image_np, 2, keepdims=True), (1, 1, 3)).astype(np.uint8)
32
33 # The input needs to be a tensor, convert it using `tf.convert_to_tensor`.
34 input_tensor = tf.convert_to_tensor(image_np)
35 # The model expects a batch of images, so add an axis with `tf.newaxis`.
36 input_tensor = input_tensor[tf.newaxis, ...]
37
38 detections = detect_fn(input_tensor)
39
40 # All outputs are batches tensors.
41 # Convert to numpy arrays, and take index [0] to remove the batch dimension.
42 # We're only interested in the first num_detections.
43 num_detections = int(detections.pop('num_detections'))
44 detections = {key: value[0, :num_detections].numpy()
45               for key, value in detections.items()}
46 detections['num_detections'] = num_detections
47
48 # detection_classes should be ints.
49 detections['detection_classes'] = detections['detection_classes'].astype(np.int64)
50
51 image_np_with_detections = image_np.copy()
52
53 viz_utils.visualize_boxes_and_labels_on_image_array(
54     image_np_with_detections,
55     detections['detection_boxes'],
56     detections['detection_classes'],
57     detections['detection_scores'],
58     category_index,
59     use_normalized_coordinates=True,
60     max_boxes_to_draw=200,
61     min_score_thresh=.4, # Adjust this value to set the minimum probability boxes to be class
62     agnostic_mode=False)
63
64 cv2.imshow('image_np_with_detections')

```





Inference Result

You can use the above script to fashion it into using a video as an input and perform inference on that.

Conclusion

Congratulations! You have built an object detection model with TensorFlow 2.

That's it for the tutorial! Hope you face no issues while following along, if there are any questions please comment and I will respond to your queries.

Refer to the [Tensorflow Github page](#).

My socials -> [LinkedIn](#), [Medium](#)

If you like this article or any of my other articles, please consider supporting me on my [Patreon](#)!

Thank you for reading!

[TensorFlow](#)[Object Detection](#)[Python](#)[Deep Learning](#)[Data Science](#)[About](#) [Help](#) [Legal](#)

Get the Medium app



